

EXECUTIVE BRIEFING

Charter

An internal AI system that reads a Business Requirements Document and drafts a scored, build-ready delivery plan — engineering plan, schedule, architecture, proof-of-concept scope, and technology options — in about ten minutes.

• Live • Azure, Central India

Single-instance internal tool

Powered by OpenAI GPT-4.1

THE PROBLEM

**Turning a BRD into a delivery plan
is slow, manual, and inconsistent**


Once a Business Requirements Document is signed off, an Engineering Manager still has to hand-produce five separate artifacts before a team can start: a phased plan, an effort and schedule estimate, a solution architecture, a proof-of-concept scope, and a technology recommendation. Each depends on the others, each should reflect the organization's own patterns and standards, and today none of that is checked before it lands on a manager's desk.

That work typically takes days, its quality depends on who happens to write it, and there is no consistent way to confirm a plan actually covers every requirement or matches how the organization has built similar things before. Charter automates the first draft of that work and grounds it in the organization's own knowledge — then scores it before a person ever sees it.

HOW IT WORKS

One BRD in, a scored delivery plan out

Charter is not one model answering a prompt — it's a coordinator, five specialists, and a reviewer, wired as a graph with one deliberate loop back into itself.



The **Orchestrator** assigns each specialist its slice of the BRD. The five specialists then run in a fixed sequence, each reading the previous one's output. The **Critic** scores every result against the organization's own bar; anything below it is sent back to the

specialist that produced it — up to two rounds — before **Assemble** compiles the final, scored document.

WHAT MAKES IT RELIABLE

Four disciplines, working together

01 Ingestion & parsing Reads the uploaded document, splits it into sections, and classifies every sentence as functional, non-functional, a constraint, or an assumption — flagging anything ambiguous.	02 Knowledge grounding Every agent retrieves relevant material from the organization's own past BRDs, templates, architecture patterns, and technology standards before it writes a word.	03 Multi-agent generation Five specialists each produce one deliverable, building on each other's output — the schedule reads the plan's phases, the PoC reads the architecture's components.	04 Validation & evaluation A Critic scores every deliverable against the organization's own bar and forces revisions — nothing reaches the Engineering Manager unreviewed.
---	---	---	--

THE ROSTER

Eight AI roles, each doing one job

Modeled on how a real planning team divides the work — a coordinator, five specialists, and a reviewer — with the model choice for each role matched to how demanding its task is.

ROLE	RESPONSIBILITY	MODEL
------	----------------	-------

BRD Ingest & Parse	Reads the document and classifies every requirement	GPT-4.1 mini
Orchestrator	Routes sections to specialists and manages the run	GPT-4.1
Engineering Plan Generator	Phases, risks, milestones, team composition — with a self-review pass before it hands off	GPT-4.1
Schedule Estimator	Effort, timeline, and resourcing aligned to the plan's phases	GPT-4.1
Solution Architect	System design, components, data flow, and non-functional coverage	GPT-4.1
PoC Planner	Proof-of-concept scope with measurable, falsifiable success criteria	GPT-4.1
Tech Stack Recommender	Two to three stack options weighed on cost, risk, and team fit	GPT-4.1
Critic	Scores every deliverable and enforces the revision loop	GPT-4.1

THE OUTPUT

Five deliverables, one document

Pl Engineering Plan Phases, risk register, milestones, team composition	Sc Schedule & Estimates Effort, timeline, resourcing, critical path

Ar

Solution Architecture

Components, data flow, NFR mapping, key decisions

Pc

PoC Plan

Scope, success criteria, exit decision

Ts

Tech Stack Options

2–3 stacks, trade-offs, a recommendation

THE QUALITY GATE

Nothing ships without a score

Every deliverable is scored on four dimensions and rolled up into one of three badges — the same scorecard the Engineering Manager sees.

DELIVERABLE	BADGE	OVERALL	COMPLETE	CONSISTENT	ACTIONABLE	GROUNDING
Engineering Plan	green	0.86	0.90	0.85	0.84	0.85
Schedule & Estimates	amber	0.72	0.78	0.66	0.75	0.70
Solution Architecture	green	0.86	0.88	0.84	0.83	0.87
PoC Plan	green	0.94	0.95	0.93	0.94	0.94
Tech Stack Options	green	0.98	0.99	0.97	0.98	0.98

Green ≥ 0.80 · Amber ≥ 0.60 · Red below 0.60. An amber or red score sends the deliverable back to its author with the Critic's specific issues attached — up to two rounds — before the run is allowed to finish.

PROOF, NOT A MOCK-UP

A real run through the system

Sample analysis — an actual internal Charter run, shown here as a concrete illustration.

~9 min end-to-end

41 Sections parsed	249 Requirements classified	5 / 5 Deliverables produced	0.93 Average quality score	Green Overall readiness
--	---	---	--	---

WHAT IT COSTS TO RUN

Infrastructure is fixed; usage is per analysis

Azure hosting

\$20–45 / month

- One always-on container, sized for internal traffic
- Registry, storage, secrets, and monitoring
- No dedicated infrastructure headcount required

OpenAI usage

\$0.50–1.50 /

analysis

- Scales with how many BRDs are analyzed, not with idle time
- Cheaper model reserved for document parsing
- Revision rounds are the only thing that raises the per-run cost

WHERE IT STANDS

Status and what's next

TODAY	Running end-to-end on Azure as a single-instance internal tool. An Engineering Manager uploads a BRD, watches (or steps away from) a background run, and reviews the scored result — with a history of every past run kept on file.
SCOPE	Built for one team's use at a time. There is no multi-user login yet — everyone who reaches the tool shares one workspace.
NEXT, IF USAGE GROWS	Individual sign-in, a managed job queue in place of the current background-run model, and expanding the knowledge base with more of the organization's own delivered projects — which directly improves how well new plans are grounded.

Charter — internal engineering tool

Built on LangGraph, retrieval-augmented generation, and OpenAI GPT-4.1